

Notebook

October 14, 2025

Flight Fare Prediction

Data Set Information: Nowadays airline tickets can vary dynmically and significantly for the same flight. customers are seeking to get the lowest prices for their flights. so here we introduces our model to save money for customers by predicting the flights fares taking various features into considerations such as flight time, destination, source, dep time , arrival time etc.. Attribute Information: Airline : names of airline companies Date_of_Journey - day/month/year Source - city from where journey starts Destination - journey ending city Route - way or direction of flight Dep_Time - the time when a flight leaves the gate(hour:minute) Arrival_Time - the time when a flight arrives the gate(hour:minute) Duration - hour:minute Total_Stops - number of stops Additional_Info - extra information Price - fare of a flight

0.1 Data Manipulation

0.2 Importing libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
[2]: pd.set_option('display.max_columns',None) #displays max number of cols
```

0.3 Importing dataset

```
[3]: train_data=pd.read_excel("Data_Train.xlsx")
```

0.4 Dataset View

```
[4]: train_data.columns
```

```
[4]: Index(['Airline', 'Date_of_Journey', 'Source', 'Destination', 'Route',
        'Dep_Time', 'Arrival_Time', 'Duration', 'Total_Stops',
        'Additional_Info', 'Price'],
        dtype='object')
```

```
[5]: train_data.head()
```

```
[5]:
```

	Airline	Date_of_Journey	Source	Destination	Route \
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR → DEL
1	Air India	1/05/2019	Kolkata	Banglore	CCU → IXR → BBI → BLR
2	Jet Airways	9/06/2019	Delhi	Cochin	DEL → LKO → BOM → COK
3	IndiGo	12/05/2019	Kolkata	Banglore	CCU → NAG → BLR
4	IndiGo	01/03/2019	Banglore	New Delhi	BLR → NAG → DEL

	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info	Price
0	22:20	01:10 22 Mar	2h 50m	non-stop	No info	3897
1	05:50	13:15	7h 25m	2 stops	No info	7662
2	09:25	04:25 10 Jun	19h	2 stops	No info	13882
3	18:05	23:30	5h 25m	1 stop	No info	6218
4	16:50	21:35	4h 45m	1 stop	No info	13302

0.5 Dataset Information

Here we can observe different datatypes like int64,object

```
[6]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10683 non-null  object
1   Date_of_Journey        10683 non-null  object
2   Source                 10683 non-null  object
3   Destination            10683 non-null  object
4   Route                  10682 non-null  object
5   Dep_Time               10683 non-null  object
6   Arrival_Time           10683 non-null  object
7   Duration               10683 non-null  object
8   Total_Stops            10682 non-null  object
9   Additional_Info        10683 non-null  object
10  Price                  10683 non-null  int64
dtypes: int64(1), object(10)
memory usage: 918.2+ KB
```

```
[7]: train_data.shape
```

```
[7]: (10683, 11)
```

0.6 Summary Statistics

Brief Information of different descriptive statistics-

Measures of Frequency :- Count, Percent, Frequency. Measures of Central Tendency :- Mean, Median, and Mode. Measures of Dispersion or Variation:- Range(min,max), Variance, Standard

Deviation. Measures of Position :- Percentile Ranks, Quartile Ranks.

```
[8]: train_data.describe()
```

```
[8]:
```

	Price
count	10683.000000
mean	9087.064121
std	4611.359167
min	1759.000000
25%	5277.000000
50%	8372.000000
75%	12373.000000
max	79512.000000

0.7 Checking for unique values in all attribute

Different numbers of distinct values in each column. Our target variable is Price.

```
[9]: train_data.nunique().sort_values(ascending=True)
```

```
[9]: Source          5
Total_Stops        5
Destination        6
Additional_Info    10
Airline            12
Date_of_Journey    44
Route             128
Dep_Time           222
Duration           368
Arrival_Time       1343
Price              1870
dtype: int64
```

0.8 Checking for missing values in each column

No such missing values in our dataset. If you want to learn how to treat the missing values. Go through this link [CLICK HERE](#)

```
[10]: !pip install missingno
import missingno as msno
msno.matrix(train_data, labels=train_data.columns, figsize=(30,16), fontsize=12)
```


```
Requirement already satisfied: missingno in
c:\users\bindunalli\anaconda3\lib\site-packages (0.5.1)
Requirement already satisfied: matplotlib in
c:\users\bindunalli\anaconda3\lib\site-packages (from missingno) (3.5.1)
Requirement already satisfied: numpy in c:\users\bindunalli\anaconda3\lib\site-
packages (from missingno) (1.21.5)
Requirement already satisfied: seaborn in
```

```

c:\users\bindunalli\anaconda3\lib\site-packages (from missingno) (0.11.2)
Requirement already satisfied: scipy in c:\users\bindunalli\anaconda3\lib\site-
packages (from missingno) (1.7.3)
Requirement already satisfied: pyparsing>=2.2.1 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(3.0.4)
Requirement already satisfied: cycler>=0.10 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(0.11.0)
Requirement already satisfied: pillow>=6.2.0 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(9.0.1)
Requirement already satisfied: packaging>=20.0 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(21.3)
Requirement already satisfied: python-dateutil>=2.7 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(2.8.2)
Requirement already satisfied: fonttools>=4.22.0 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
c:\users\bindunalli\anaconda3\lib\site-packages (from matplotlib->missingno)
(1.3.2)
Requirement already satisfied: six>=1.5 in
c:\users\bindunalli\anaconda3\lib\site-packages (from python-
dateutil>=2.7->matplotlib->missingno) (1.16.0)
Requirement already satisfied: pandas>=0.23 in
c:\users\bindunalli\anaconda3\lib\site-packages (from seaborn->missingno)
(1.4.2)
Requirement already satisfied: pytz>=2020.1 in
c:\users\bindunalli\anaconda3\lib\site-packages (from
pandas>=0.23->seaborn->missingno) (2021.3)

```

[10]: <AxesSubplot:>



```
[11]: train_data.isnull().sum()
```

```
[11]: Airline          0
      Date_of_Journey  0
      Source          0
      Destination     0
      Route           1
      Dep_Time        0
      Arrival_Time    0
      Duration        0
      Total_Stops     1
      Additional_Info  0
      Price           0
      dtype: int64
```


```
[12]: train_data.dropna(inplace= True)    #dropping Nan values
      train_data.isnull().sum()
```

```
[12]: Airline          0
      Date_of_Journey  0
      Source          0
      Destination     0
      Route           0
      Dep_Time        0
      Arrival_Time    0
      Duration        0
      Total_Stops     0
```


```
Additional_Info    0
Price             0
dtype: int64
```

1 Analysing Categorical Variables


```
[13]: Airline_var=pd.crosstab(index=train_data['Airline'],columns='% observations')
plt.pie(Airline_var['% observations'],labels=Airline_var['% observations'].
        ↪index,autopct='%0f%%',radius=1.4)
plt.title('Airlines')
plt.show()
```



```
[14]: Source_var=pd.crosstab(index=train_data['Source'],columns='% observations')
plt.pie(Source_var['% observations'],labels=Source_var['% observations'].
        ↪index,autopct='%1.1f%%',radius=0.8)
plt.title('Source ')
plt.show()
```



```
[15]: Destination_var=pd.crosstab(index=train_data['Destination'],columns='%_
↳observations')
plt.pie(Destination_var['% observations'],labels=Destination_var['%_
↳observations'].index,autopct='%1.1f%%',radius=0.8)
plt.title('Destination')
plt.show()
```



2 EDA

From description we can see that Date_of_Journey is a object data type, *Therefore, we have to convert this datatype into timestamp so as to use this column properly for prediction* For this we require pandas to_datetime to convert object data type to datetime dtype .dt.day method will extract only day of that date .dt.month method will extract only month of that date

```
[16]: #here date and time is of string so to use them we will convert them into date,
      ↪time type and use to_datetime()
      # .dt.date method will extract only date
      # .dt.month will extract the month

train_data["Journey_day"] = pd.to_datetime(train_data.Date_of_Journey,
      ↪format="%d/%m/%Y").dt.day
train_data["Journey_month"] = pd.to_datetime(train_data.Date_of_Journey,
      ↪format="%d/%m/%Y").dt.month
#train_data["Journey_year"] = pd.to_datetime(train_data.Date_of_Journey,
      ↪format="%d/%m/%Y").dt.year

# Since we have converted Date_of_Journey column into integers, Now we can drop
      ↪as it is of no use.

train_data.drop(["Date_of_Journey"], axis = 1, inplace = True)
```

```
[17]: train_data.head()
```

```
[17]:
```

	Airline	Source	Destination	Route	Dep_Time	\
0	IndiGo	Banglore	New Delhi	BLR → DEL	22:20	
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	
2	Jet Airways	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	
3	IndiGo	Kolkata	Banglore	CCU → NAG → BLR	18:05	
4	IndiGo	Banglore	New Delhi	BLR → NAG → DEL	16:50	

	Arrival_Time	Duration	Total_Stops	Additional_Info	Price	Journey_day	\
0	01:10 22 Mar	2h 50m	non-stop	No info	3897	24	
1	13:15	7h 25m	2 stops	No info	7662	1	
2	04:25 10 Jun	19h	2 stops	No info	13882	9	
3	23:30	5h 25m	1 stop	No info	6218	12	
4	21:35	4h 45m	1 stop	No info	13302	1	

	Journey_month
0	3
1	5
2	6


```
3          5
4          3
```

```
[18]: #similarly we will extract minute and seconds from dep_time()
train_data["Dep_hour"] = pd.to_datetime(train_data["Dep_Time"]).dt.hour
train_data["Dep_min"] = pd.to_datetime(train_data["Dep_Time"]).dt.minute

# Now we can drop Dep_Time as it is of no use
train_data.drop(["Dep_Time"], axis = 1, inplace = True)
```

```
[19]: train_data.head()
```

```
[19]:
```

	Airline	Source	Destination	Route	Arrival_Time	\
0	IndiGo	Banglore	New Delhi	BLR → DEL	01:10 22 Mar	
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	13:15	
2	Jet Airways	Delhi	Cochin	DEL → LKO → BOM → COK	04:25 10 Jun	
3	IndiGo	Kolkata	Banglore	CCU → NAG → BLR	23:30	
4	IndiGo	Banglore	New Delhi	BLR → NAG → DEL	21:35	

	Duration	Total_Stops	Additional_Info	Price	Journey_day	Journey_month	\
0	2h 50m	non-stop	No info	3897	24	3	
1	7h 25m	2 stops	No info	7662	1	5	
2	19h	2 stops	No info	13882	9	6	
3	5h 25m	1 stop	No info	6218	12	5	
4	4h 45m	1 stop	No info	13302	1	3	

	Dep_hour	Dep_min
0	22	20
1	5	50
2	9	25
3	18	5
4	16	50

```
[20]: # Similar to Date_of_Journey we can extract values from dt.hour() and dt.min()

train_data["Arrival_hour"] = pd.to_datetime(train_data.Arrival_Time).dt.hour
train_data["Arrival_min"] = pd.to_datetime(train_data.Arrival_Time).dt.minute

# Now we can drop Arrival_Time as it is of no use
train_data.drop(["Arrival_Time"], axis = 1, inplace = True)
```

```
[21]: train_data.head()
```

```
[21]:
```

	Airline	Source	Destination	Route	Duration	\
0	IndiGo	Banglore	New Delhi	BLR → DEL	2h 50m	
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	7h 25m	
2	Jet Airways	Delhi	Cochin	DEL → LKO → BOM → COK	19h	

3	IndiGo	Kolkata	Bangalore	CCU → NAG → BLR	5h 25m
4	IndiGo	Bangalore	New Delhi	BLR → NAG → DEL	4h 45m

	Total_Stops	Additional_Info	Price	Journey_day	Journey_month	Dep_hour	\
0	non-stop	No info	3897	24	3	22	
1	2 stops	No info	7662	1	5	5	
2	2 stops	No info	13882	9	6	9	
3	1 stop	No info	6218	12	5	18	
4	1 stop	No info	13302	1	3	16	

	Dep_min	Arrival_hour	Arrival_min
0	20	1	10
1	50	13	15
2	25	4	25
3	5	23	30
4	50	21	35

```
[22]: # Assigning and converting Duration column into list
duration = list(train_data["Duration"])

for i in range(len(duration)):
    if len(duration[i].split()) != 2:    # Check if duration contains only hour
        or mins
        if "h" in duration[i]:
            duration[i] = duration[i].strip() + " 0m"    # Adds 0 minute
        else:
            duration[i] = "0h " + duration[i]              # Adds 0 hour

duration_hours = []
duration_mins = []
for i in range(len(duration)):
    duration_hours.append(int(duration[i].split(sep = "h")[0]))    # Extract
    or hours from duration
    duration_mins.append(int(duration[i].split(sep = "m")[0].split()[-1]))    #
    or Extracts only minutes from duration
```

```
[23]: # Adding duration_hours and duration_mins list to train_data dataframe

train_data["Duration_hours"] = duration_hours
train_data["Duration_mins"] = duration_mins
```

```
[24]: #dropping Duration
train_data.drop(["Duration"], axis = 1, inplace = True)
train_data.head()
```

	Airline	Source	Destination	Route	Total_Stops	\
0	IndiGo	Bangalore	New Delhi	BLR → DEL	non-stop	

1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	2 stops
2	Jet Airways	Delhi	Cochin	DEL → LKO → BOM → COK	2 stops
3	IndiGo	Kolkata	Banglore	CCU → NAG → BLR	1 stop
4	IndiGo	Banglore	New Delhi	BLR → NAG → DEL	1 stop

	Additional_Info	Price	Journey_day	Journey_month	Dep_hour	Dep_min	\
0	No info	3897	24	3	22	20	
1	No info	7662	1	5	5	50	
2	No info	13882	9	6	9	25	
3	No info	6218	12	5	18	5	
4	No info	13302	1	3	16	50	

	Arrival_hour	Arrival_min	Duration_hours	Duration_mins
0	1	10	2	50
1	13	15	7	25
2	4	25	19	0
3	23	30	5	25
4	21	35	4	45

3 Handling Categorical Data

One can find many ways to handle categorical data. Some of them categorical data are, *Nominal data* → data are not in any order → *OneHotEncoder* is used in this case Ordinal data → data are in order → *LabelEncoder* is used in this case

```
[25]: train_data["Airline"].value_counts()
```


```
[25]: Jet Airways          3849
      IndiGo              2053
      Air India           1751
      Multiple carriers    1196
      SpiceJet             818
      Vistara              479
      Air Asia            319
      GoAir                194
      Multiple carriers Premium economy  13
      Jet Airways Business    6
      Vistara Premium economy    3
      Trujet                1
      Name: Airline, dtype: int64
```

```
[26]: # From graph we can see that Jet Airways Business have the highest Price.
      # Apart from the first Airline almost all are having similar median

      # Airline vs Price
      sns.set(rc={"figure.figsize":(30,10)})
      sns.boxplot(y=train_data["Price"], x = train_data["Airline"])
```

```
plt.show()
```

#Inference: Here with the help of the cat plot we are trying to plot the
↪ boxplot between the price of the flight and airline
#and we can conclude that Jet Airways has the most outliers in terms of price.



[27]: *# As Airline is Nominal Categorical data we will perform OneHotEncoding*

```
Airline = train_data[["Airline"]]

Airline = pd.get_dummies(Airline, drop_first= True)

Airline.head()
```

```
[27]:   Airline_Air India  Airline_GoAir  Airline_IndiGo  Airline_Jet Airways \
0                0                0                1                0
1                1                0                0                0
2                0                0                0                1
3                0                0                1                0
4                0                0                1                0

   Airline_Jet Airways Business  Airline_Multiple carriers \
0                        0                        0
1                        0                        0
2                        0                        0
3                        0                        0
4                        0                        0

   Airline_Multiple carriers Premium economy  Airline_SpiceJet \
0                        0                        0
1                        0                        0
2                        0                        0
3                        0                        0
4                        0                        0
```


	Airline_Trujet	Airline_Vistara	Airline_Vistara Premium economy
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

```
[28]: train_data["Source"].value_counts()
```

```
[28]: Delhi      4536
      Kolkata    2871
      Bangalore  2197
      Mumbai     697
      Chennai    381
      Name: Source, dtype: int64
```

```
[29]: # Source vs Price
```

```
sns.set(rc={"figure.figsize":(30,10)})
sns.boxplot(y = train_data["Price"], x = train_data["Source"])
plt.show()
```



```
[30]: # As Source is Nominal Categorical data we will perform OneHotEncoding
```

```
Source = train_data[["Source"]]

Source = pd.get_dummies(Source, drop_first= True)

Source.head()
```

```
[30]: Source_Chennai Source_Delhi Source_Kolkata Source_Mumbai
      0              0              0              0
```

1	0	0	1	0
2	0	1	0	0
3	0	0	1	0
4	0	0	0	0

```
[31]: train_data["Destination"].value_counts()
```

```
[31]: Cochin      4536
Banglore    2871
Delhi       1265
New Delhi   932
Hyderabad   697
Kolkata     381
Name: Destination, dtype: int64
```

```
[32]: # As Destination is Nominal Categorical data we will perform OneHotEncoding

Destination = train_data[["Destination"]]

Destination = pd.get_dummies(Destination, drop_first = True)

Destination.head()
```

```
[32]: Destination_Cochin  Destination_Delhi  Destination_Hyderabad  \
0                0                0                0
1                0                0                0
2                1                0                0
3                0                0                0
4                0                0                0

Destination_Kolkata  Destination_New Delhi
0                0                1
1                0                0
2                0                0
3                0                0
4                0                1
```

```
[33]: train_data.head()
```

```
[33]: Airline  Source Destination  Route Total_Stops  \
0      IndiGo  Banglore   New Delhi    BLR → DEL    non-stop
1    Air India  Kolkata    Banglore  CCU → IXR → BBI → BLR    2 stops
2  Jet Airways   Delhi    Cochin    DEL → LKO → BOM → COK    2 stops
3      IndiGo  Kolkata    Banglore    CCU → NAG → BLR    1 stop
4      IndiGo  Banglore   New Delhi    BLR → NAG → DEL    1 stop

Additional_Info  Price  Journey_day  Journey_month  Dep_hour  Dep_min  \
```

0	No info	3897	24	3	22	20
1	No info	7662	1	5	5	50
2	No info	13882	9	6	9	25
3	No info	6218	12	5	18	5
4	No info	13302	1	3	16	50

	Arrival_hour	Arrival_min	Duration_hours	Duration_mins
0	1	10	2	50
1	13	15	7	25
2	4	25	19	0
3	23	30	5	25
4	21	35	4	45

```
[34]: # Additional_Info contains almost 80% no_info
# Route and Total_Stops are related to each other so drop Route and use
      ↳ Total_stops

train_data.drop(["Route", "Additional_Info"], axis = 1, inplace = True)
      ↳ #dropping column of missing values since it is of no use
train_data["Total_Stops"].value_counts()
```

```
[34]: 1 stop      5625
non-stop    3491
2 stops     1520
3 stops      45
4 stops       1
Name: Total_Stops, dtype: int64
```

```
[35]: # As this is case of Ordinal Categorical type we perform LabelEncoder
# Here Values are assigned with corresponding keys

train_data.replace({"non-stop": 0, "1 stop": 1, "2 stops": 2, "3 stops": 3, "4
      ↳ stops": 4}, inplace = True)
train_data.head()
```

```
[35]:      Airline  Source Destination  Total_Stops  Price  Journey_day \
0      IndiGo  Bangalore  New Delhi          0   3897          24
1    Air India  Kolkata    Bangalore          2   7662           1
2  Jet Airways    Delhi    Cochin           2  13882           9
3      IndiGo  Kolkata    Bangalore          1   6218          12
4      IndiGo  Bangalore  New Delhi          1  13302           1

      Journey_month  Dep_hour  Dep_min  Arrival_hour  Arrival_min \
0                3         22        20            1          10
1                5          5         50           13          15
2                6          9         25            4          25
3                5         18          5           23          30
```

4	3	16	50	21	35
---	---	----	----	----	----

	Duration_hours	Duration_mins
0	2	50
1	7	25
2	19	0
3	5	25
4	4	45

```
[36]: # Concatenate dataframe --> train_data + Airline + Source + Destination

train_data1 = pd.concat([train_data, Airline, Source, Destination], axis = 1)
#concatenating column-wise
train_data1.head()
```

```
[36]:
```

	Airline	Source	Destination	Total_Stops	Price	Journey_day	\
0	IndiGo	Banglore	New Delhi	0	3897	24	
1	Air India	Kolkata	Banglore	2	7662	1	
2	Jet Airways	Delhi	Cochin	2	13882	9	
3	IndiGo	Kolkata	Banglore	1	6218	12	
4	IndiGo	Banglore	New Delhi	1	13302	1	

	Journey_month	Dep_hour	Dep_min	Arrival_hour	Arrival_min	\
0	3	22	20	1	10	
1	5	5	50	13	15	
2	6	9	25	4	25	
3	5	18	5	23	30	
4	3	16	50	21	35	

	Duration_hours	Duration_mins	Airline_Air India	Airline_GoAir	\
0	2	50	0	0	
1	7	25	1	0	
2	19	0	0	0	
3	5	25	0	0	
4	4	45	0	0	

	Airline_IndiGo	Airline_Jet Airways	Airline_Jet Airways Business	\
0	1	0	0	
1	0	0	0	
2	0	1	0	
3	1	0	0	
4	1	0	0	

	Airline_Multiple carriers	Airline_Multiple carriers Premium economy	\
0	0	0	
1	0	0	
2	0	0	

3	0	0
4	0	0

	Airline_SpiceJet	Airline_Trujet	Airline_Vistara \
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

	Airline_Vistara Premium economy	Source_Chennai	Source_Delhi \
0	0	0	0
1	0	0	0
2	0	0	1
3	0	0	0
4	0	0	0

	Source_Kolkata	Source_Mumbai	Destination_Cochin	Destination_Delhi \
0	0	0	0	0
1	1	0	0	0
2	0	0	1	0
3	1	0	0	0
4	0	0	0	0

	Destination_Hyderabad	Destination_Kolkata	Destination_New Delhi
0	0	0	1
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	1

```
[37]: train_data1.drop(["Airline", "Source", "Destination"], axis = 1, inplace = True)
train_data1.head()
```

```
[37]:
```

	Total_Stops	Price	Journey_day	Journey_month	Dep_hour	Dep_min \
0	0	3897	24	3	22	20
1	2	7662	1	5	5	50
2	2	13882	9	6	9	25
3	1	6218	12	5	18	5
4	1	13302	1	3	16	50

	Arrival_hour	Arrival_min	Duration_hours	Duration_mins \
0	1	10	2	50
1	13	15	7	25
2	4	25	19	0
3	23	30	5	25
4	21	35	4	45

	Airline_Air India	Airline_GoAir	Airline_IndiGo	Airline_Jet Airways	\
0	0	0	1	0	
1	1	0	0	0	
2	0	0	0	1	
3	0	0	1	0	
4	0	0	1	0	

	Airline_Jet Airways Business	Airline_Multiple carriers	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	Airline_Multiple carriers Premium economy	Airline_SpiceJet	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	Airline_Trujet	Airline_Vistara	Airline_Vistara Premium economy	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	Source_Chennai	Source_Delhi	Source_Kolkata	Source_Mumbai	\
0	0	0	0	0	
1	0	0	1	0	
2	0	1	0	0	
3	0	0	1	0	
4	0	0	0	0	

	Destination_Cochin	Destination_Delhi	Destination_Hyderabad	\
0	0	0	0	
1	0	0	0	
2	1	0	0	
3	0	0	0	
4	0	0	0	

	Destination_Kolkata	Destination_New Delhi
0	0	1
1	0	0
2	0	0

3	0	0
4	0	1

```
[38]: train_data1.shape
```

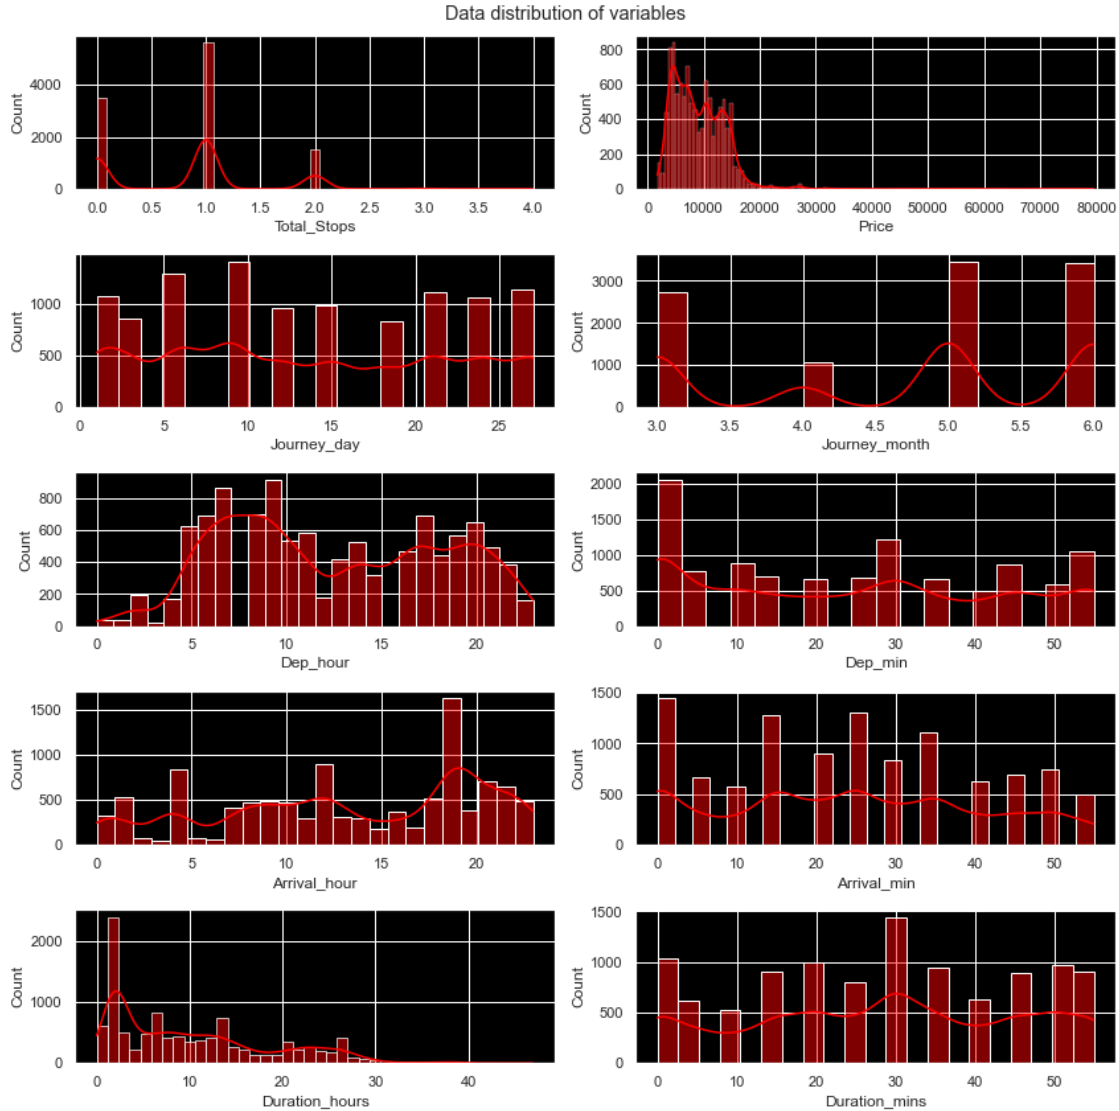
```
[38]: (10682, 30)
```

3.1 Checking the data distribution of each variable

Skewed Distribution-

What is skewed distribution? If one tail is longer than another, the distribution is skewed. These distributions are sometimes called asymmetric or asymmetrical distributions as they don't show any kind of symmetry. Symmetry means that one half of the distribution is a mirror image of the other half. For example, the normal distribution is a symmetric distribution with no skew. The tails are exactly the same. Left Skewed or Negatively Skewed:- A left-skewed distribution has a long left tail. Left-skewed distributions are also called negatively-skewed distributions. (Mean < Median < Mode) Right Skewed or Positively Skewed:- A right-skewed distribution has a long right tail. Right-skewed distributions are also called positive-skew distributions. (Mean > Median > Mode) Symmetric Distribution:- A symmetric distribution is a type of distribution where the left side of the distribution mirrors the right side (Mean = Median = Mode). ex- Normal Distribution

```
[39]: plt.figure(figsize=(12, 12))
      for i, col in enumerate(train_data1.select_dtypes(include=['float', 'int64']).
      ↪columns):
          plt.rcParams['axes.facecolor'] = 'black'
          ax = plt.subplot(5, 2, i+1)
          sns.histplot(data=train_data1, x=col, ax=ax, color='red', kde=True)
      plt.suptitle('Data distribution of variables')
      plt.tight_layout()
```




4 Box plot(Outliers Detection)

Box Plot-

What is Box Plot? In descriptive statistics, a box plot or boxplot is a method for graphically demonstrating the locality, spread and skewness groups of numerical data through their quartiles.

How to interpret boxplot *Median: In the box plot, the median is displayed rather than the mean.*
 Q1: The first quartile (25%) position. * Q3: The third quartile (75%) position. * Interquartile range (IQR): a measure of statistical dispersion, being equal to the difference between 75th and 25th percentiles. It represents how 50% of the points were dispersed. * Lower and upper $1.5IQR$ whiskers: *These represent the limits and boundaries for the outliers.* Outliers: Defined as observations that fall below $Q1 - 1.5 IQR$ or above $Q3 + 1.5 IQR$. Outliers are displayed as dots or circles.

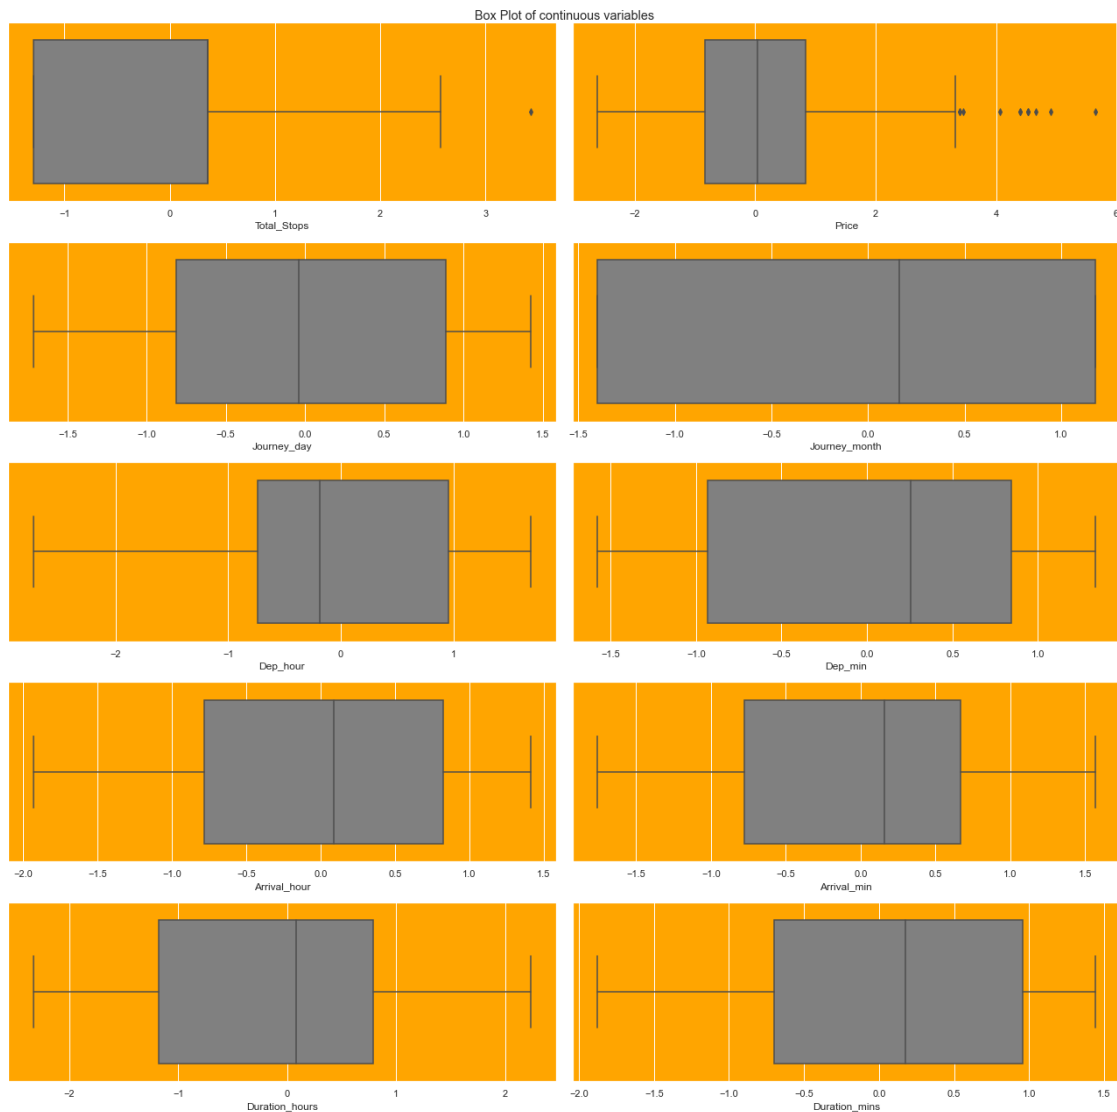
```
[40]: plt.figure(figsize=(18, 18))
for i, col in enumerate(train_data1.select_dtypes(include=['float64','int64']).
    columns):
    plt.rcParams['axes.facecolor'] = 'White'
    ax = plt.subplot(5,2, i+1)
    sns.boxplot(data=train_data1, x=col, ax=ax,color='Purple')
plt.suptitle('Box Plot of continuous variables')
plt.tight_layout()
```



4.1 Data distribution after applying Power Transformer

```
[41]: #selecting variables that have data types float and int.
var=list(train_data1.select_dtypes(include=['float64','int64']).columns)
from sklearn.preprocessing import PowerTransformer
sc_X=PowerTransformer(method = 'yeo-johnson')
train_data1[var]=sc_X.fit_transform(train_data1[var])

[42]: plt.figure(figsize=(18, 18))
for i, col in enumerate(train_data1.select_dtypes(include=['float64','int64']).
    ↪columns):
    plt.rcParams['axes.facecolor'] = 'Orange'
    ax = plt.subplot(5,2, i+1)
    sns.boxplot(data=train_data1, x=col, ax=ax,color='Grey')
plt.suptitle('Box Plot of continuous variables')
plt.tight_layout()
```



5 Feature Selection

Feature Selection-

Feature selection methods are intended to reduce the number of input variables to those that are believed to be most useful to a model in order to predict the target variable..

In our dataset we have numerical Input variable and numerical Output variable.so we will use correlation for the feature selection.

```
[43]: train_data1.shape
```

```
[43]: (10682, 30)
```

```
[44]: train_data1.columns
```

```
[44]: Index(['Total_Stops', 'Price', 'Journey_day', 'Journey_month', 'Dep_hour',  
        'Dep_min', 'Arrival_hour', 'Arrival_min', 'Duration_hours',  
        'Duration_mins', 'Airline_Air India', 'Airline_GoAir', 'Airline_IndiGo',  
        'Airline_Jet Airways', 'Airline_Jet Airways Business',  
        'Airline_Multiple carriers',  
        'Airline_Multiple carriers Premium economy', 'Airline_SpiceJet',  
        'Airline_Trujet', 'Airline_Vistara', 'Airline_Vistara Premium economy',  
        'Source_Chennai', 'Source_Delhi', 'Source_Kolkata', 'Source_Mumbai',  
        'Destination_Cochin', 'Destination_Delhi', 'Destination_Hyderabad',  
        'Destination_Kolkata', 'Destination_New Delhi'],  
        dtype='object')
```

```
[45]: X = train_data1.loc[:, ['Total_Stops', 'Journey_day', 'Journey_month',  
        ↪ 'Dep_hour',  
        'Dep_min', 'Arrival_hour', 'Arrival_min', 'Duration_hours',  
        'Duration_mins', 'Airline_Air India', 'Airline_GoAir', 'Airline_IndiGo',  
        'Airline_Jet Airways', 'Airline_Jet Airways Business',  
        'Airline_Multiple carriers',  
        'Airline_Multiple carriers Premium economy', 'Airline_SpiceJet',  
        'Airline_Trujet', 'Airline_Vistara', 'Airline_Vistara Premium economy',  
        'Source_Chennai', 'Source_Delhi', 'Source_Kolkata', 'Source_Mumbai',  
        'Destination_Cochin', 'Destination_Delhi', 'Destination_Hyderabad',  
        'Destination_Kolkata', 'Destination_New Delhi']]  
X.head()
```

```
[45]:
```

	Total_Stops	Journey_day	Journey_month	Dep_hour	Dep_min	Arrival_hour	\
0	-1.297820	1.164296	-1.401748	1.545888	0.023186	-1.790733	
1	1.574617	-1.716424	0.161418	-1.356237	1.179354	-0.056006	
2	1.574617	-0.405463	1.175096	-0.548198	0.255935	-1.362584	

3	0.358782	-0.041621	0.161418	0.956329	-0.933677	1.413910
4	0.358782	-1.716424	-1.401748	0.646652	1.179354	1.118899

	Arrival_min	Duration_hours	Duration_mins	Airline_Air	India	\
0	-0.776578	-1.175643	1.200413		0	
1	-0.433010	-0.055254	-0.099976		1	
2	0.156840	1.074715	-1.877928		0	
3	0.420855	-0.393117	-0.099976		0	
4	0.670321	-0.603213	0.955571		0	

	Airline_GoAir	Airline_IndiGo	Airline_Jet Airways	\
0	0	1	0	
1	0	0	0	
2	0	0	1	
3	0	1	0	
4	0	1	0	

	Airline_Jet Airways Business	Airline_Multiple carriers	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	Airline_Multiple carriers Premium economy	Airline_SpiceJet	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	Airline_Trujet	Airline_Vistara	Airline_Vistara Premium economy	\
0	0	0		0
1	0	0		0
2	0	0		0
3	0	0		0
4	0	0		0

	Source_Chennai	Source_Delhi	Source_Kolkata	Source_Mumbai	\
0	0	0	0	0	
1	0	0	1	0	
2	0	1	0	0	
3	0	0	1	0	
4	0	0	0	0	

	Destination_Cochin	Destination_Delhi	Destination_Hyderabad	\
0	0	0	0	

1	0	0	0
2	1	0	0
3	0	0	0
4	0	0	0

	Destination_Kolkata	Destination_New Delhi
0	0	1
1	0	0
2	0	0
3	0	0
4	0	1

```
[46]: y = train_data1.iloc[:, 1]
      y.head()
```

```
[46]: 0    -1.367854
      1    -0.138984
      2     1.086164
      3    -0.536300
      4     0.993291
      Name: Price, dtype: float64
```

```
[47]: from sklearn.feature_selection import SelectKBest
      from sklearn.feature_selection import r_regression    #Correlation
      rs = SelectKBest(score_func=r_regression, k='all')
      rs.fit(X, y)
```

```
[47]: SelectKBest(k='all', score_func=<function r_regression at 0x00000195F100FE50>)
```

```
[48]: feature_contribution=(rs.scores_/sum(rs.scores_))*100
```


```
[49]: for i,j in enumerate(X.columns):
      print(f'{j} : {feature_contribution[i]:.2f}%')
```

```
Total_Stops : 217.98%
Journey_day : -40.15%
Journey_month : -18.16%
Dep_hour : 1.83%
Dep_min : -21.69%
Arrival_hour : 16.16%
Arrival_min : -32.82%
Duration_hours : 216.12%
Duration_mins : -44.70%
Airline_Air India : 23.27%
Airline_GoAir : -31.81%
Airline_IndiGo : -120.31%
Airline_Jet Airways : 137.46%
Airline_Jet Airways Business : 34.22%
```

Airline_Multiple carriers : 52.89%
Airline_Multiple carriers Premium economy : 7.01%
Airline_SpiceJet : -113.96%
Airline_Trujet : -3.77%
Airline_Vistara : -13.87%
Airline_Vistara Premium economy : 0.57%
Source_Chennai : -65.89%
Source_Delhi : 105.00%
Source_Kolkata : 11.52%
Source_Mumbai : -100.02%
Destination_Cochin : 105.00%
Destination_Delhi : -105.05%
Destination_Hyderabad : -100.02%
Destination_Kolkata : -65.89%
Destination_New Delhi : 49.10%

[50]: *# Finds correlation between Independent and dependent attributes*

```
plt.figure(figsize = (8,8))  
sns.heatmap(train_data.corr(),annot = True, cmap = "RdYlGn")  
  
plt.show()
```



```
[51]: # Important feature selection using ExtraTreesRegressor
```

```
from sklearn.ensemble import ExtraTreesRegressor
selection = ExtraTreesRegressor()
selection.fit(X, y)
```

```
[51]: ExtraTreesRegressor()
```


```
[52]: #plot graph of feature importances for better visualization
```

```
plt.figure(figsize = (12,8))
feat_importances = pd.Series(selection.feature_importances_, index=X.columns)
```

```

feat_importances.nlargest(20).plot(kind='barh')
plt.ylabel("Feature labels")
plt.xlabel("Feature Importance")
plt.show()

```



```

[53]: X1 = train_data1.loc[:,['Total_Stops', 'Journey_day', 'Journey_month',
    ↳ 'Dep_hour',
    ↳ 'Dep_min', 'Arrival_hour', 'Arrival_min', 'Duration_hours',
    ↳ 'Duration_mins', 'Airline_IndiGo',
    ↳ 'Airline_Jet Airways','Airline_Air India',
    ↳ 'Airline_Multiple carriers', 'Source_Delhi',
    ↳ 'Destination_Cochin', 'Destination_New
    ↳ Delhi', 'Source_Mumbai', 'Destination_Hyderabad', 'Airline_SpiceJet', 'Airline_Jet
    ↳ Airways Business'],]

```

5.0.1 Splitting our dataset into train and test set

```

[54]: #splitting our dataset in 80% training and 20% testset
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X1, y, test_size = 0.2,
    ↳ random_state = 42)

```

5.0.2 Feature Scaling

Feature Scaling-

What is Normalization? Normalization is a scaling technique in which values are shifted and

rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling. What is Standardization? Standardization is another scaling technique where the values are centered around the mean with a unit standard deviation. This means that the mean of the attribute becomes zero and the resultant distribution has a unit standard deviation.

Here we are going to use Standardization.

</html>

```
[55]: from sklearn.preprocessing import StandardScaler
      sc=StandardScaler()
      X_train=sc.fit_transform(X_train)
      X_test=sc.transform(X_test)
```

6 Fitting model

Fitting model 1. Split dataset into train and test set in order to prediction w.r.t X_{test} 2.If needed do scaling of data 3.Scaling is not done in Random forest 4.Import model 5.Fit the data 6.Predict w.r.t X_{test} 7.In regression check RSME Score 8.Plot graph

```
[56]: from sklearn.ensemble import RandomForestRegressor
      from sklearn.linear_model import LinearRegression
      from sklearn import metrics
      from sklearn.metrics import r2_score
      from sklearn.metrics import mean_squared_error
      from sklearn.metrics import mean_absolute_error
      from math import sqrt
      from sklearn.model_selection import GridSearchCV
      from sklearn.model_selection import RandomizedSearchCV
      import statsmodels
```

```
[57]: #creating dictionary for storing different models accuracy
      model_comparison={}
```

```
[58]: import statsmodels.api as sm
      lr=sm.OLS(y_train,X_train).fit()
      print(lr.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:          Price    R-squared (uncentered):
0.707
Model:                OLS      Adj. R-squared (uncentered):
0.706
Method:              Least Squares    F-statistic:
1142.
Date:                Tue, 29 Nov 2022    Prob (F-statistic):
0.00
```

Time: 13:48:59 Log-Likelihood:
-6850.7
No. Observations: 8545 AIC:
1.374e+04
Df Residuals: 8527 BIC:
1.386e+04
Df Model: 18
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
x1	0.4162	0.012	34.853	0.000	0.393	0.440
x2	-0.1143	0.006	-19.322	0.000	-0.126	-0.103
x3	-0.0883	0.007	-13.345	0.000	-0.101	-0.075
x4	0.0140	0.006	2.343	0.019	0.002	0.026
x5	-0.0343	0.006	-5.701	0.000	-0.046	-0.023
x6	-0.0014	0.006	-0.225	0.822	-0.013	0.011
x7	0.0011	0.006	0.174	0.862	-0.011	0.013
x8	0.0995	0.012	8.161	0.000	0.076	0.123
x9	-0.0075	0.006	-1.226	0.220	-0.020	0.005
x10	-0.0762	0.009	-8.123	0.000	-0.095	-0.058
x11	0.3217	0.011	29.566	0.000	0.300	0.343
x12	0.0497	0.009	5.281	0.000	0.031	0.068
x13	0.1666	0.009	18.513	0.000	0.149	0.184
x14	0.0112	0.004	2.811	0.005	0.003	0.019
x15	0.0112	0.004	2.811	0.005	0.003	0.019
x16	0.0944	0.007	13.664	0.000	0.081	0.108
x17	-0.0604	0.003	-18.508	0.000	-0.067	-0.054
x18	-0.0604	0.003	-18.508	0.000	-0.067	-0.054
x19	-0.1107	0.008	-14.320	0.000	-0.126	-0.096
x20	0.0997	0.006	16.957	0.000	0.088	0.111

Omnibus: 562.652 Durbin-Watson: 1.991
Prob(Omnibus): 0.000 Jarque-Bera (JB): 1811.427
Skew: 0.305 Prob(JB): 0.00
Kurtosis: 5.172 Cond. No. 2.24e+16

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The smallest eigenvalue is 6.39e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

7 Linear Regression

```
[59]: model=LinearRegression()  
      model.fit(X_train, y_train)  
      y_pred= model.predict(X_test)
```

```
[60]: model.score(X_test, y_test)
```

```
[60]: 0.7045632169213696
```

```
[61]: model.score(X_train, y_train)
```

```
[61]: 0.7068898950577607
```

```
[62]: metrics.r2_score(y_test, y_pred)
```

```
[62]: 0.7045632169213696
```

```
[63]: print('MAE:', metrics.mean_absolute_error(y_test,y_pred))  
      print('MSE:', metrics.mean_squared_error(y_test,y_pred))  
      print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test,y_pred)))  
      model_comparison['Linear_  
      ↪Regression']=[r2_score(y_test,y_pred),mean_squared_error(y_test,y_pred),mean_absolute_error_  
      ↪sqrt(metrics.mean_squared_error(y_test,y_pred))]
```

```
MAE: 0.4106656782564049
```

```
MSE: 0.3039384913943007
```

```
RMSE: 0.5513061684711144
```

8 Checking linearity assumption

```
[64]: def calculate_residuals(model, features, label):  
      # Creates predictions on the features with the model and calculates residuals  
      predictions = model.predict(features)  
      df_results = pd.DataFrame({'Actual': label, 'Predicted': predictions})  
      df_results['Residuals'] = abs(df_results['Actual']) -_  
      ↪abs(df_results['Predicted'])  
  
      return df_results
```

```
[65]: def linear_assumption(model, features, label):  
      # Linearity: Assumes that there is a linear relationship between the predictors_  
      ↪and the response variable.  
      #If not, either a quadratic term or another algorithm should be used.  
      print('Assumption 1: Linear Relationship between the Target and the_  
      ↪Feature', '\n')
```

```

print('Checking with a scatter plot of actual vs. predicted.',
      'Predictions should follow the diagonal line.')

# Calculating residuals for the plot
df_results = calculate_residuals(model, features, label)

# Plotting the actual vs predicted values
sns.lmplot(x='Actual', y='Predicted', data=df_results, fit_reg=False,
↪size=7)

# Plotting the diagonal line
line_coords = np.arange(df_results.min().min(), df_results.max().max())
plt.plot(line_coords, line_coords, # X and y points
         color='darkorange', linestyle='--')
plt.title('Actual vs. Predicted')
plt.show()

```

```
[66]: linear_assumption(model,X_train,y_train)
```


Assumption 1: Linear Relationship between the Target and the Feature

Checking with a scatter plot of actual vs. predicted. Predictions should follow the diagonal line.

```

C:\Users\Bindunalli\anaconda3\lib\site-packages\seaborn\regression.py:581:
UserWarning: The `size` parameter has been renamed to `height`; please update
your code.
  warnings.warn(msg, UserWarning)

```

9 normality assumption

```
[67]: def normal_errors_assumption(model, features, label, p_value_thresh=0.05):
#Normality: Assumes that the error terms are normally distributed.
#If they are not, nonlinear transformations of variables may solve this.
#This assumption being violated primarily causes issues with the confidence_
↪intervals
    from statsmodels.stats.diagnostic import normal_ad
    print('Assumption 2: The error terms are normally distributed', '\n')

    # Calculating residuals for the Anderson-Darling test
```

```

df_results = calculate_residuals(model, features, label)

print('Using the Anderson-Darling test for normal distribution')

# Performing the test on the residuals
p_value = normal_ad(df_results['Residuals'])[1]
print('p-value from the test - below 0.05 generally means non-normal:',
↪p_value)

# Reporting the normality of the residuals
if p_value < p_value_thresh:
    print('Residuals are not normally distributed')
else:
    print('Residuals are normally distributed')

# Plotting the residuals distribution
plt.subplots(figsize=(12, 6))
plt.title('Distribution of Residuals')
sns.distplot(df_results['Residuals'])
plt.show()

print()
if p_value > p_value_thresh:
    print('Assumption satisfied')
else:
    print('Assumption not satisfied')
    print()
    print('Confidence intervals will likely be affected')
    print('Try performing nonlinear transformations on variables')

```

```
[68]: normal_errors_assumption(model,X_train,y_train)
```

Assumption 2: The error terms are normally distributed

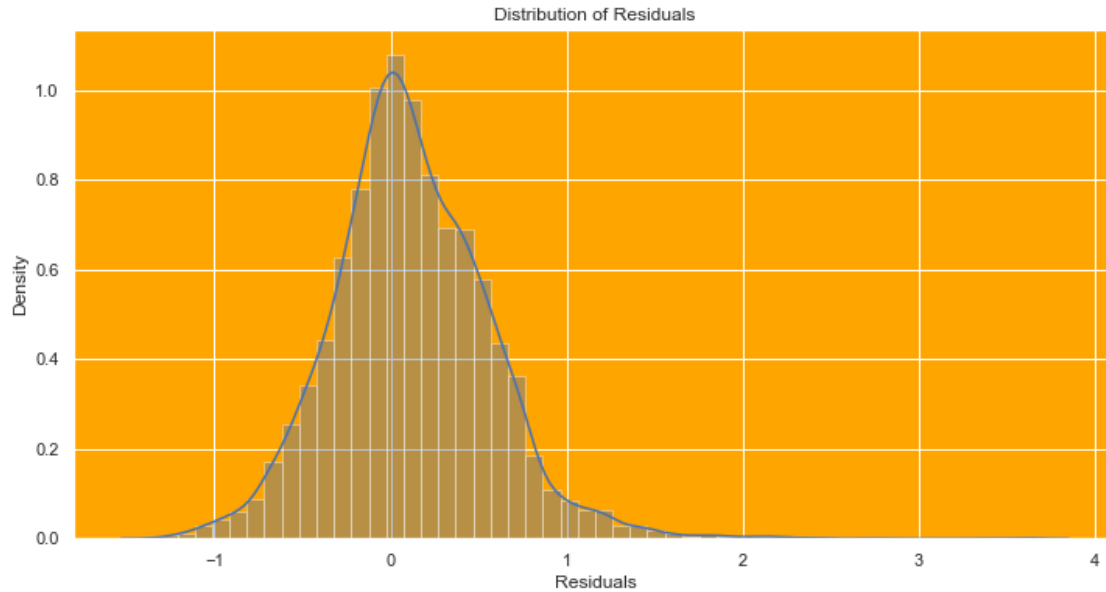
Using the Anderson-Darling test for normal distribution

p-value from the test - below 0.05 generally means non-normal: 0.0

Residuals are not normally distributed

C:\Users\Bindunalli\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).

```
warnings.warn(msg, FutureWarning)
```



Assumption not satisfied

Confidence intervals will likely be affected

Try performing nonlinear transformations on variables

10 multicollinearity assumption

```
[69]: from statsmodels.stats.outliers_influence import variance_inflation_factor

def calc_vif(X):

    # Calculating VIF
    vif = pd.DataFrame()
    vif["variables"] = X.columns
    vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.
↪shape[1])]

    return(vif)
```

```
[70]: calc_vif(X.select_dtypes(include=['float', 'int64']))
```

```
[70]:
```

	variables	VIF
0	Total_Stops	3.651543
1	Journey_day	1.002697
2	Journey_month	1.022830
3	Dep_hour	1.018446

```

4      Dep_min    1.021049
5      Arrival_hour 1.055272
6      Arrival_min 1.089885
7      Duration_hours 3.676177
8      Duration_mins 1.077575

```

11 autocorrelation assumption

```

[71]: def autocorrelation_assumption(model, features, label):
      """
      Autocorrelation: Assumes that there is no autocorrelation in the residuals.
      ↪If there is autocorrelation, then there is a pattern that is not
      ↪explained due to the current value being dependent on the previous value.
      This may be resolved by adding a lag variable of either
      ↪the dependent variable or some of the predictors.
      """
      from statsmodels.stats.stattools import durbin_watson
      print('Assumption 4: No Autocorrelation', '\n')

      # Calculating residuals for the Durbin Watson-tests
      df_results = calculate_residuals(model, features, label)

      print('\nPerforming Durbin-Watson Test')
      print('Values of 1.5 < d < 2.5 generally show that there is no
      ↪autocorrelation in the data')
      print('0 to 2< is positive autocorrelation')
      print('>2 to 4 is negative autocorrelation')
      print('-----')
      durbinWatson = durbin_watson(df_results['Residuals'])
      print('Durbin-Watson:', durbinWatson)
      if durbinWatson < 1.5:
          print('Signs of positive autocorrelation', '\n')
          print('Assumption not satisfied')
      elif durbinWatson > 2.5:
          print('Signs of negative autocorrelation', '\n')
          print('Assumption not satisfied')
      else:
          print('Little to no autocorrelation', '\n')
          print('Assumption satisfied')

```

```

[72]: autocorrelation_assumption(model,X_train,y_train)

```

Assumption 4: No Autocorrelation

Performing Durbin-Watson Test
 Values of $1.5 < d < 2.5$ generally show that there is no autocorrelation in the data
 0 to 2< is positive autocorrelation
 >2 to 4 is negative autocorrelation

 Durbin-Watson: 1.8623720423867665
 Little to no autocorrelation

 Assumption satisfied

12 Homoscedasticity assumption

```
[73]: def homoscedasticity_assumption(model, features, label):
        """
        Homoscedasticity: Assumes that the errors exhibit constant variance
        """
        print('Assumption 5: Homoscedasticity of Error Terms', '\n')

        print('Residuals should have relative constant variance')

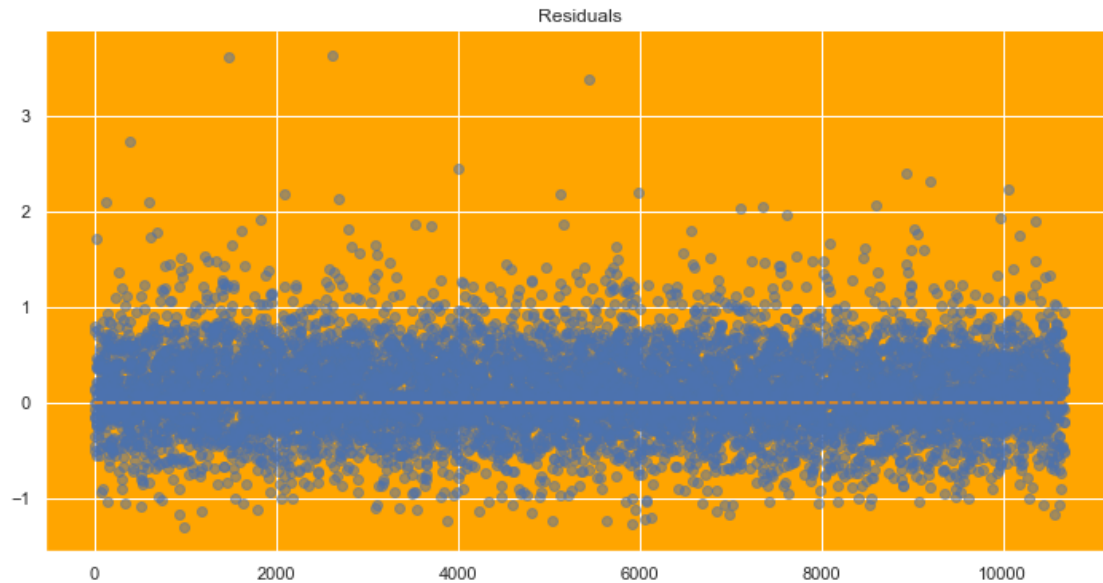
        # Calculating residuals for the plot
        df_results = calculate_residuals(model, features, label)

        # Plotting the residuals
        plt.subplots(figsize=(12, 6))
        ax = plt.subplot(111) # To remove spines
        plt.scatter(x=df_results.index, y=df_results.Residuals, alpha=0.5)
        plt.plot(np.repeat(0, df_results.index.max()), color='darkorange',
        ↵linestyle='--')
        ax.spines['right'].set_visible(False) # Removing the right spine
        ax.spines['top'].set_visible(False) # Removing the top spine
        plt.title('Residuals')
        plt.show()
```

```
[74]: homoscedasticity_assumption(model,X_train,y_train)
```

Assumption 5: Homoscedasticity of Error Terms

Residuals should have relative constant variance



13 Decision Tree Regression model

```
[75]: # Training the Decision Tree Regression model
from sklearn.tree import DecisionTreeRegressor
regressor = DecisionTreeRegressor(random_state = 0)
regressor.fit(X_train,y_train)
```

```
[75]: DecisionTreeRegressor(random_state=0)
```

```
[76]: # Predicting test set results
y_pred = regressor.predict(X_test)
print('Train Score:',regressor.score(X_train,y_train))
print('Test Score:',regressor.score(X_test,y_test))
print(y_pred,y_test)
```

```
Train Score: 0.9706998735408532
```

```
Test Score: 0.7834265202708964
```

```
[ 1.51647721 -0.94857957  0.05348141 ... -0.2659763  0.99931154
 1.15606124] 6075      1.491422
```

```
3544      -0.948580
```

```
9291       0.220000
```

```
5032     -1.384920
```

```
2483       0.926663
```

```
...
```

```
9797     -0.204246
```

```
9871     -1.073097
```

```
10063    -0.192814
```

```
8802      0.139151
8617      1.128185
Name: Price, Length: 2137, dtype: float64
```

```
[77]: from sklearn.metrics import r2_score
      r2_score(y_test, y_pred)
```

```
[77]: 0.7834265202708964
```

```
[78]: print('MAE:', metrics.mean_absolute_error(y_test,y_pred))
      print('MSE:', metrics.mean_squared_error(y_test,y_pred))
      print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
      model_comparison['Decision Tree_
      ↳Regression']=[r2_score(y_test,y_pred),mean_squared_error(y_test,y_pred),mean_absolute_error
      ↳sqrt(metrics.mean_squared_error(y_test,y_pred))]
```

```
MAE: 0.2850029639750224
MSE: 0.22280575904916547
RMSE: 0.4720230492774325
```

13.1 RANDOM FOREST (linearity assumption is violated so random forest(non-linear data) is used).

```
[79]: reg_rf = RandomForestRegressor()
      reg_rf.fit(X_train, y_train)
      y_pred = reg_rf.predict(X_test)
```

```
[80]: reg_rf.score(X_train, y_train)
```

```
[80]: 0.9602301178291667
```

```
[81]: reg_rf.score(X_test, y_test)
```

```
[81]: 0.8620813490497017
```

```
[82]: metrics.r2_score(y_test, y_pred)
```

```
[82]: 0.8620813490497017
```


```
[83]: print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
      print('MSE:', metrics.mean_squared_error(y_test, y_pred))
      print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
      model_comparison['Random Forest_
      ↳Regression']=[r2_score(y_test,y_pred),mean_squared_error(y_test,y_pred),mean_absolute_error
      ↳sqrt(metrics.mean_squared_error(y_test,y_pred))]
```

```
MAE: 0.24807512855846778
MSE: 0.14188750049385054
RMSE: 0.37667957270583513
```


```
[84]: sns.distplot(y_test-y_pred)
plt.show()
```

C:\Users\Bindunalli\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).

```
warnings.warn(msg, FutureWarning)
```



```
[85]: plt.scatter(y_test, y_pred, alpha = 0.5)
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.show()
```



14 Hyperparameter Tuning

1. Choose following method for hyperparameter tuning
2. RandomizedSearchCV -> Fast

3. GridSearchCV
4. Assign hyperparameters in form of dictionary
5. Fit the model
6. Check best parameters and best score

```
[86]: from sklearn.model_selection import RandomizedSearchCV
      #Randomized Search CV

      # Number of trees in random forest
      n_estimators = [int(x) for x in np.linspace(start = 100, stop = 1200, num = 12)]
      # Number of features to consider at every split
      max_features = ['auto', 'sqrt']
      # Maximum number of levels in tree
      max_depth = [int(x) for x in np.linspace(5, 30, num = 6)]
      # Minimum number of samples required to split a node
      min_samples_split = [2, 5, 10, 15, 100]
      # Minimum number of samples required at each leaf node
      min_samples_leaf = [1, 2, 5, 10]
```

```
[87]: # Create the random grid

      random_grid = {'n_estimators': n_estimators,
                     'max_features': max_features,
                     'max_depth': max_depth,
                     'min_samples_split': min_samples_split,
                     'min_samples_leaf': min_samples_leaf}
```

```
[88]: # Random search of parameters, using 5 fold cross validation,
      # search across 100 different combinations
      rf_random = RandomizedSearchCV(estimator = reg_rf, param_distributions =_
      ↪random_grid,scoring='neg_mean_squared_error',
                                     n_iter = 10, cv = 5, verbose=2)
```

```
[89]: rf_random.fit(X_train,y_train)
```

```
Fitting 5 folds for each of 10 candidates, totalling 50 fits
[CV] END max_depth=15, max_features=auto, min_samples_leaf=10,
min_samples_split=2, n_estimators=700; total time= 7.8s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=10,
min_samples_split=2, n_estimators=700; total time= 7.9s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=10,
min_samples_split=2, n_estimators=700; total time= 8.1s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=10,
min_samples_split=2, n_estimators=700; total time= 7.9s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=10,
min_samples_split=2, n_estimators=700; total time= 7.9s
[CV] END max_depth=30, max_features=auto, min_samples_leaf=5,
min_samples_split=2, n_estimators=800; total time= 10.4s
```



```

[CV] END max_depth=10, max_features=auto, min_samples_leaf=5,
min_samples_split=100, n_estimators=200; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=5,
min_samples_split=100, n_estimators=200; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=5,
min_samples_split=100, n_estimators=200; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=5,
min_samples_split=100, n_estimators=200; total time= 1.7s
[CV] END max_depth=10, max_features=auto, min_samples_leaf=5,
min_samples_split=100, n_estimators=200; total time= 1.7s
[CV] END max_depth=25, max_features=sqrt, min_samples_leaf=1,
min_samples_split=100, n_estimators=700; total time= 2.2s
[CV] END max_depth=25, max_features=sqrt, min_samples_leaf=1,
min_samples_split=100, n_estimators=700; total time= 2.2s
[CV] END max_depth=25, max_features=sqrt, min_samples_leaf=1,
min_samples_split=100, n_estimators=700; total time= 2.2s
[CV] END max_depth=25, max_features=sqrt, min_samples_leaf=1,
min_samples_split=100, n_estimators=700; total time= 2.1s
[CV] END max_depth=25, max_features=sqrt, min_samples_leaf=1,
min_samples_split=100, n_estimators=700; total time= 2.2s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=5,
min_samples_split=10, n_estimators=800; total time= 3.3s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=5,
min_samples_split=10, n_estimators=800; total time= 3.4s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=5,
min_samples_split=10, n_estimators=800; total time= 3.3s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=5,
min_samples_split=10, n_estimators=800; total time= 3.3s
[CV] END max_depth=20, max_features=sqrt, min_samples_leaf=5,
min_samples_split=10, n_estimators=800; total time= 3.4s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=1,
min_samples_split=10, n_estimators=400; total time= 5.4s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=1,
min_samples_split=10, n_estimators=400; total time= 5.3s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=1,
min_samples_split=10, n_estimators=400; total time= 5.2s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=1,
min_samples_split=10, n_estimators=400; total time= 5.2s
[CV] END max_depth=15, max_features=auto, min_samples_leaf=1,
min_samples_split=10, n_estimators=400; total time= 5.3s

```

```

[89]: RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(),
                        param_distributions={'max_depth': [5, 10, 15, 20, 25, 30],
                                             'max_features': ['auto', 'sqrt'],
                                             'min_samples_leaf': [1, 2, 5, 10],
                                             'min_samples_split': [2, 5, 10, 15,
                                                                    100]},

```

```
        'n_estimators': [100, 200, 300, 400,
                          500, 600, 700, 800,
                          900, 1000, 1100,
                          1200]},
        scoring='neg_mean_squared_error', verbose=2)
```

```
[90]: rf_random.best_params_
```

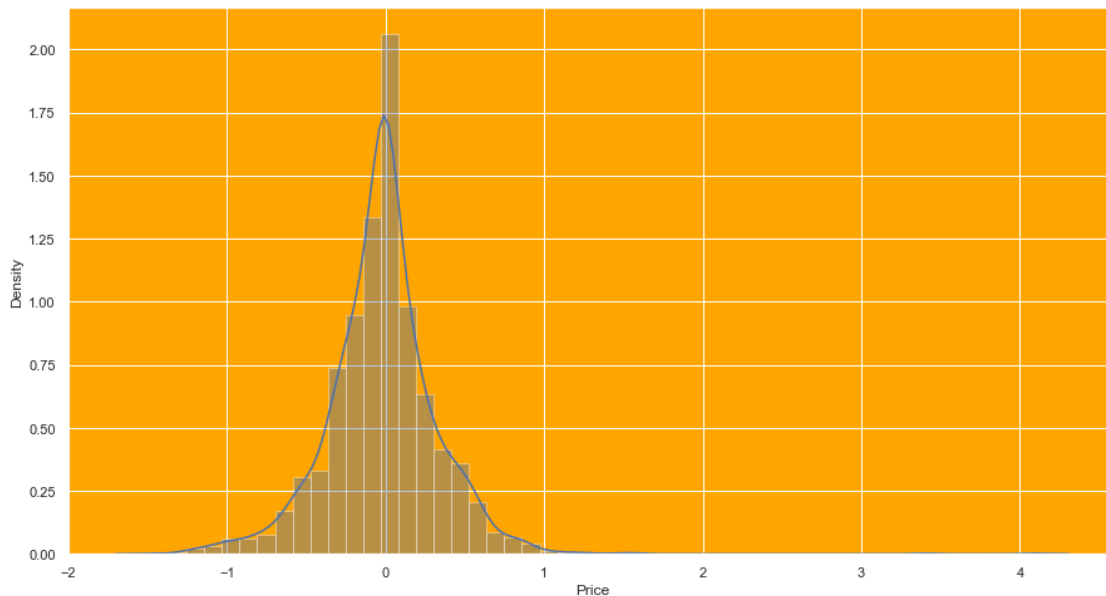
```
[90]: {'n_estimators': 400,
      'min_samples_split': 10,
      'min_samples_leaf': 1,
      'max_features': 'auto',
      'max_depth': 15}
```

```
[91]: prediction = rf_random.predict(X_test)
```


```
[92]: plt.figure(figsize = (15,8))
      sns.distplot(y_test-prediction)
      plt.show()
```

C:\Users\Bindunalli\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).

```
warnings.warn(msg, FutureWarning)
```



```
[93]: plt.figure(figsize = (8,8))
plt.scatter(y_test, prediction, alpha = 0.5)
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.show()
```



```
[94]: print('MAE:', metrics.mean_absolute_error(y_test, prediction))
print('MSE:', metrics.mean_squared_error(y_test, prediction))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, prediction)))
model_comparison['After_
↳hypertuning']=[r2_score(y_test,prediction),mean_squared_error(y_test,prediction),mean_absol
↳sqrt(metrics.mean_squared_error(y_test,prediction))]
```

MAE: 0.24142034750191704

MSE: 0.124074760458832
RMSE: 0.3522424739562678

```
[95]: metrics.r2_score(y_test,prediction)
```

```
[95]: 0.879395834587944
```

15 Model Comparison

```
[96]: Model_com_df=pd.DataFrame(model_comparison).T  
Model_com_df.columns=['R-Square','MSE','MAE','RMSE']  
Model_com_df=Model_com_df.sort_values(by='R-Square',ascending=False)  
Model_com_df.style.format("{:.2%}").background_gradient(cmap='Blues')
```

```
[96]: <pandas.io.formats.style.Styler at 0x19580a11f10>
```

This notebook was converted with [convert.ploomber.io](https://ploomber.io)